

AI-Powered Personalized Learning Chatbot Using Large Language Models



Aicha Rikli, Ragad Alrowili, Kholod Almutairi, Lena Bin Ghannam.


Supervised by:
Dr. Nourah Alangari

Abstract:

This research addresses the challenges of adapting Large Language Models (LLMs) for educational tutoring by fine-tuning Qwen2.5-14B using Quantized Low-Rank Adaptation (QLoRA) and Group Relative Policy Optimization (GRPO). We explored the effects of varying LoRA ranks and prompting strategies on mathematical reasoning. Results indicate that lower LoRA ranks (rank=8) yield superior performance with minimal policy drift (achieving 91.66% Pass@1 accuracy on the GSM8K Dataset), while few-shot prompting significantly enhances zero-shot accuracy. However, a significant pedagogical-reasoning gap emerged: enforcing Socratic constraints at inference time reduced raw mathematical accuracy by approximately 7%. To assess instructional quality, we evaluated the model using a pedagogical rubric assessed by human experts and an LLM-as-a-Judge (GPT-4o). Ultimately, these findings highlight the critical trade-off between raw logical precision and instructional utility in educational AI.

Methodology:

To address the computational constraints of training large language models, we selected the Qwen2.5-14B-Instruct model and optimized it using QLoRA and the Unsloth framework for efficient fine-tuning. To specifically enhance mathematical reasoning, we employed GRPO as our reinforcement learning method. Our experimental design involved a systematic evaluation of architectural and prompting variables, testing different LoRA ranks (8, 32, 128) and prompting strategies (zero-shot vs. few-shot) to explore their impact on reasoning accuracy. Following the identification of the best-performing LoRA rank, we investigated the pedagogical-reasoning trade-off by applying a Socratic pedagogical prompt at inference time to two distinct configurations: the best-performing reasoning model (LoRA Rank 8) and a specialized LoRA Rank 8 model fine-tuned with a pedagogical system prompt. Performance was evaluated quantitatively via Pass@1 accuracy and qualitatively using a pedagogical rubric assessed by human experts (n=3) and an LLM-as-a-Judge (GPT-4o).



Results:

Our experiments identified LoRA Rank 8 as the most efficient configuration, yielding the highest zero-shot accuracy (91.66%) and the lowest KL divergence (Table 1). While 4-shot prompting effectively boosted performance across the base and Rank 8 models. However, a pedagogical prompt at inference time caused accuracy to plummet to 84.38% for the reasoning-tuned model and 83.62% for the pedagogical-tuned model (Table 2). Finally, Qualitative evaluation revealed a major discrepancy between evaluators, as human evaluators were much stricter and provided lower scores compared to the more lenient LLM-as-a-judge (Figure 1).

Model Configuration	zero-Shot Accuracy (%)	4-Shot Accuracy (%)	KL Divergence
Base Model (14.7B)	83.47	92.72	-
LoRA Rank 8	91.66	93.40	0.0012
LoRA Rank 32	90.98	-	0.0054
LoRA Rank 128	83.78	-	1.1289

Table 1: GSM8K Test Accuracy by LoRA Rank and Prompting Strategy

Model Configuration	Accuracy (%)
LoRA Rank 8 (Reasoning-Tuned)	84.38
LoRA Rank 8 (Pedagogical-Tuned)	83.62

Table 2: GSM8K Test Accuracy Under Inference-Time Pedagogical Prompts




Figure 1: Pedagogical Evaluation Comparison: LLM-as-a-Judge vs. Human Experts (N=11)

Conclusion:

This study demonstrates that low-rank adaptation and few-shot prompting effectively maximize mathematical accuracy without high computational costs. Additionally it confirms a critical trade-off: enforcing pedagogical constraints causes accuracy to plummet. Furthermore, our findings reiterate that human evaluation remains indispensable for detecting nuanced instructional flaws. Ultimately, developing effective AI tutors requires balancing efficient fine-tuning with essential human oversight.

References:

Please scan the QR code to view the list of references.



Graduation Project Fair 2026

College of Computer and Information Sciences